

Proseminar: High Performance Computing  
Assignment 02

## 1 Exercise 1

1.1 Write a sequential application `pi_seq` in C or C++ that computes  $\pi$  for a given number of samples (command line argument). Test your application for various, large sample sizes to verify the correctness of your implementation.

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <math.h>
4
5 double calculate_pi(int num_points) {
6     int points_in_circle = 0;
7
8     for (int count = 0; count < num_points; count++) {
9         double random_x = (double)rand() / RAND_MAX;
10        double random_y = (double)rand() / RAND_MAX;
11
12        if (random_x * random_x + random_y * random_y <= 1.0) {
13            points_in_circle++;
14        }
15    }
16
17    return 4.0 * points_in_circle / num_points;
18 }
19
20 int main(int argc, char *argv[]) {
21     int total_points = atoi(argv[1]);
22
23     double pi_approx = calculate_pi(total_points);
24
25     printf("Approximation of Pi: %f\n", pi_approx);
26
27     return 0;
28 }
```

Listing 1: Sequential Pi Approximation

1.2 Consider a parallelization strategy using MPI. Which communication pattern(s) would you choose and why?

As there are no dependencies while generating random points and checking if they lie in the circle, we can do that individually per rank. So we can just split the problem size  $N$  into  $N \div n_{ranks}$  chunks. At the end one core have to retrieve the number of points in circle from all other cores, build a sum out of it and calculate  $\pi$  with it.

For this we would use `MPI_Reduce(...)` as it can collect a variable from all ranks and directly use operations as addition by giving reduce the operation type `MPI_SUM`.

1.3 Implement your chosen parallelization strategy as a second application `pi_mpi`. Run it with varying numbers of ranks and sample sizes and verify its correctness by comparing the output to `pi_seq`.

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <math.h>
4 #include <mpi.h>
5
```

```

6 void calculate_pi(int num_points, int rank, int numProcs) {
7     // WTime Initialization
8     double starttime, endtime;
9     starttime = MPI_Wtime();
10    // Define how many points each rank generates and generate them
11    int points_for_rank = round(num_points / numProcs);
12    int points_in_circle_in_rank = 0;
13    int points_in_circle_global = 0;
14
15    for (int count = 0; count < points_for_rank; count++) {
16        double random_x = (double)rand() / RAND_MAX;
17        double random_y = (double)rand() / RAND_MAX;
18
19        if (random_x * random_x + random_y * random_y <= 1.0) {
20            points_in_circle_in_rank++;
21        }
22    }
23    // Gather results from all ranks
24    MPI_Reduce(&points_in_circle_in_rank, &points_in_circle_global, 1, MPI_INT, MPI_SUM, 0,
25    MPI_COMM_WORLD);
26
27    if (rank == 0) {
28        double pi_approx = 4.0 * points_in_circle_global / num_points;
29        endtime = MPI_Wtime();
30        printf("WTime: %f seconds\n", endtime - starttime);
31        printf("Approximation of Pi: %f\n", pi_approx);
32    }
33 }
34
35 int main(int argc, char *argv[]) {
36     // MPI Initialization
37     MPI_Init(&argc, &argv);
38     int rank, numProcs;
39     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
40     MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
41
42     int total_points = atoi(argv[1]);
43
44     if (rank == 0) {
45         printf("Total Points used for Approximation: %i \n", total_points);
46     }
47
48     calculate_pi(total_points, rank, numProcs);
49
50     MPI_Finalize();
51     return 0;
52 }

```

Listing 2: Parallel Pi Approximation

## 1.4 Discuss the effects and implications of your parallelization.

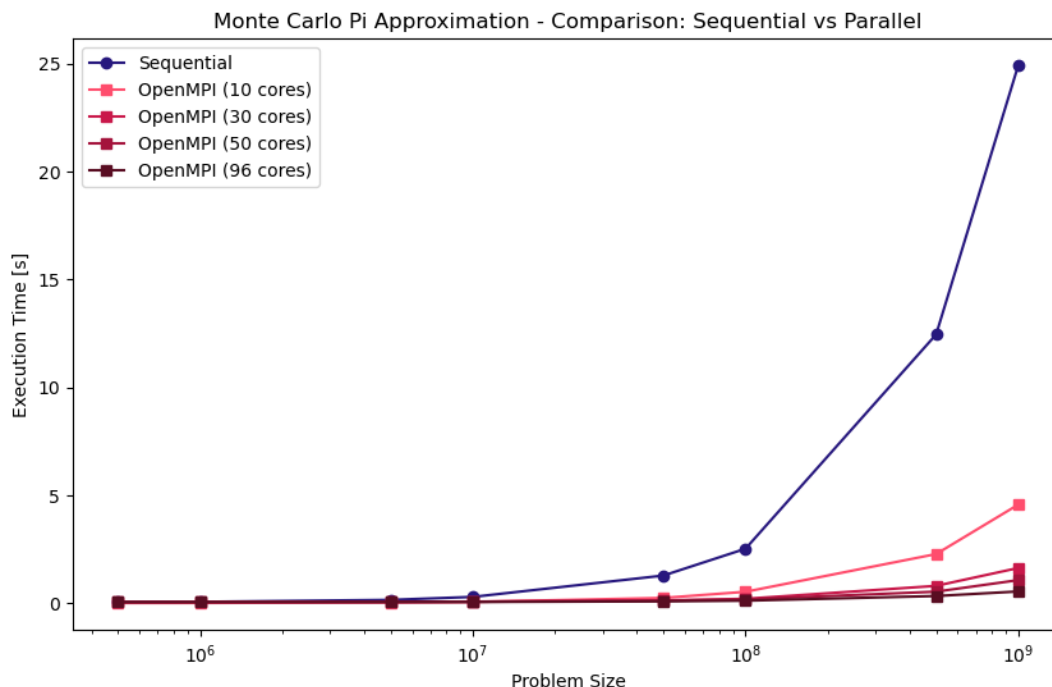


Figure 1: Monte Carlo Pi Approximation - Runtime Comparison

When running the experiment with a problem size below  $10^7$  the sequential code and the parallelized code have around the same runtimes. If more than  $10^7$  points are used the execution time of the sequential program grows rapidly. By adding more cores the highest problem size ( $10^9$ ) gets faster computed. The execution time for  $10^9$  points is 5x faster with 96 cores instead of just 10.

## 2 Exercise 2

**2.1 A sequential implementation of a 1-D heat stencil is available in `heat_stencil_1D_seq.c`. Read the code and make sure you understand what happens. See the Wikipedia article on Stencil Codes for more information.**

Iterative Stencil Loops

**2.2 Consider a parallelization strategy using MPI. Which communication pattern(s) would you choose and why? Are there additional changes required in the code beyond calling MPI functions? If so, elaborate!**

- **Boundary-Exchange-Pattern** (using `MPI_Sendrecv`), because exchanging boundary values between processes is essential for correct calculations in each section.
- **Broadcast-Pattern** (with `MPI_Bcast`), to efficiently distribute initial data to all processes.
- **Gather-Pattern** (with `MPI_Gather`), to collect results from individual processes at the end of the computation.

Additional changes that need to be made include implementing functionality that splits the problem into chunks that can be distributed via MPI. This also includes additional error handling to ensure that the input fits our strategy.

## 2.3 Implement your chosen parallelization strategy as a second application `heat_stencil_1D_mpi`. Run it with varying numbers of ranks and problem sizes and verify its correctness by comparing the output to `heat_stencil_1D_seq`.

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <mpi.h>
4
5 typedef double value_t;
6
7 #define RESOLUTION 120
8
9 // -- vector utilities --
10
11 typedef value_t *Vector;
12
13 Vector createVector(int N);
14 void releaseVector(Vector m);
15 void printTemperature(Vector m, int N);
16
17 // -- simulation code ---
18
19 int main(int argc, char **argv) {
20     // 'parsing' optional input parameter = problem size
21     int N = 2000;
22     if (argc > 1) {
23         N = atoi(argv[1]);
24     }
25     int T = N * 500;
26
27     // MPI Initialization
28     MPI_Init(&argc, &argv);
29     int rank, numProcs;
30     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
31     MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
32
33     // Time measurement variables
34     double starttime, endtime;
35
36     // Calculate chunk size and handle remainder
37     int chunkSize = N / numProcs;
38     int remainder = N % numProcs;
39
40     // Determine start and end index for each rank
41     int startIdx = rank * chunkSize + (rank < remainder ? rank : remainder);
42     int endIdx = startIdx + chunkSize + (rank < remainder ? 1 : 0);
43     int localSize = endIdx - startIdx;
44
45     // Debug information
46     printf("[DEBUG] Rank: %i, localSize: %i, startIdx: %i, endIdx: %i \n", rank, localSize, startIdx,
47           endIdx);
48
49     // Create a buffer for storing temperature fields
50     Vector A = createVector(N);
51
52     // Counts and displacements for MPI_Gatherv
53     int *counts = NULL;
54     int *displs = NULL;
55     if (rank == 0) {
56         counts = malloc(sizeof(int) * numProcs);
57         displs = malloc(sizeof(int) * numProcs);
58         for (int i = 0; i < numProcs; i++) {
59             int tempStartIdx = i * chunkSize + (i < remainder ? i : remainder);
60             int tempEndIdx = tempStartIdx + chunkSize + (i < remainder ? 1 : 0);
61             counts[i] = tempEndIdx - tempStartIdx;
62             displs[i] = tempStartIdx;
63         }
64     }
```

```

63 }
64
65 int source_x = N / 4;
66
67 if (rank == 0) {
68     printf("Computing heat-distribution for room size N=%d for T=%d timesteps\n", N, T);
69
70     // ----- setup -----
71     // Set up initial conditions in A
72     for (int i = 0; i < N; i++) {
73         A[i] = 273; // temperature is 0 C everywhere (273 K)
74     }
75
76     // Heat source in one corner
77     A[source_x] = 273 + 60;
78
79     printf("Initial:\t");
80     printTemperature(A, N);
81     printf("\n");
82 }
83
84 // Distribute A to all processes
85 MPI_Bcast(A, N, MPI_DOUBLE, 0, MPI_COMM_WORLD);
86 printf("[DEBUG] Broadcasted A to every rank for computation. \n");
87
88 // Start time measurement before the computation begins
89 starttime = MPI_Wtime();
90
91 // ----- compute -----
92
93 // Create buffers for the computation and initialize with A
94 Vector B = createVector(localSize);
95 Vector C = createVector(localSize);
96 for (int i = startIdx; i < endIdx; i++) {
97     B[i - startIdx] = A[i];
98 }
99
100 // Create variables for the neighbours
101 value_t rightNeighbour = 0;
102 value_t leftNeighbour = 0;
103
104 // For each time step
105 for (int t = 0; t < T; t++) {
106     // Exchange boundary values with neighboring processes
107     if (rank > 0) {
108         MPI_Sendrecv(&B[0], 1, MPI_DOUBLE, rank - 1, 0, &leftNeighbour, 1, MPI_DOUBLE, rank - 1, 1,
109             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
110     } else {
111         leftNeighbour = B[0];
112     }
113     if (rank < numProcs - 1) {
114         MPI_Sendrecv(&B[localSize - 1], 1, MPI_DOUBLE, rank + 1, 1, &rightNeighbour, 1, MPI_DOUBLE,
115             rank + 1, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
116     } else {
117         rightNeighbour = B[localSize - 1];
118     }
119
120     // Propagate the temperature
121     for (int i = startIdx; i < endIdx; i++) {
122         // Heat source remains constant
123         if (i == source_x) {
124             C[i - startIdx] = B[i - startIdx];
125             continue;
126         }
127
128         // Current temperature
129         value_t tc = B[i - startIdx];

```

```

129 // Temperatures of adjacent cells
130 value_t t1 = (i == startIdx) ? leftNeighbour : B[i - 1 - startIdx];
131 value_t tr = (i == endIdx - 1) ? rightNeighbour : B[i - startIdx + 1];
132
133 // Compute new temperature
134 C[i - startIdx] = tc + 0.2 * (t1 + tr + (-2 * tc));
135 }
136
137 // Swap buffers
138 Vector H = B;
139 B = C;
140 C = H;
141
142 // Print temperature at intervals
143 if (t % 10000 == 0) {
144     MPI_Gatherv(B, localSize, MPI_DOUBLE, A, counts, displs, MPI_DOUBLE, 0, MPI_COMM_WORLD);
145
146     if (rank == 0) {
147         printf("Step t=%d:\t", t);
148         printTemperature(A, N);
149         printf("\n");
150     }
151 }
152 }
153
154 // Gather final results from all processes
155 MPI_Gatherv(B, localSize, MPI_DOUBLE, A, counts, displs, MPI_DOUBLE, 0, MPI_COMM_WORLD);
156
157 releaseVector(B);
158 releaseVector(C);
159
160 // End time measurement after computation
161 endtime = MPI_Wtime();
162
163 // ----- check -----
164 int success = 1;
165
166 if (rank == 0) {
167     printf("Final:\t\t");
168     printTemperature(A, N);
169     printf("\n");
170
171     for (int i = 0; i < N; i++) {
172         value_t temp = A[i];
173         if (273 <= temp && temp <= 273 + 60)
174             continue;
175         success = 0;
176         break;
177     }
178
179     printf("Verification: %s\n", (success) ? "OK" : "FAILED");
180
181     // Print the time
182     printf("WTime: %f seconds\n", endtime - starttime);
183
184     // Free counts and displacements arrays
185     free(counts);
186     free(displs);
187 }
188
189 // ----- cleanup -----
190
191 releaseVector(A);
192
193 // Finalize MPI
194 MPI_Finalize();
195 return (success) ? EXIT_SUCCESS : EXIT_FAILURE;
196 }

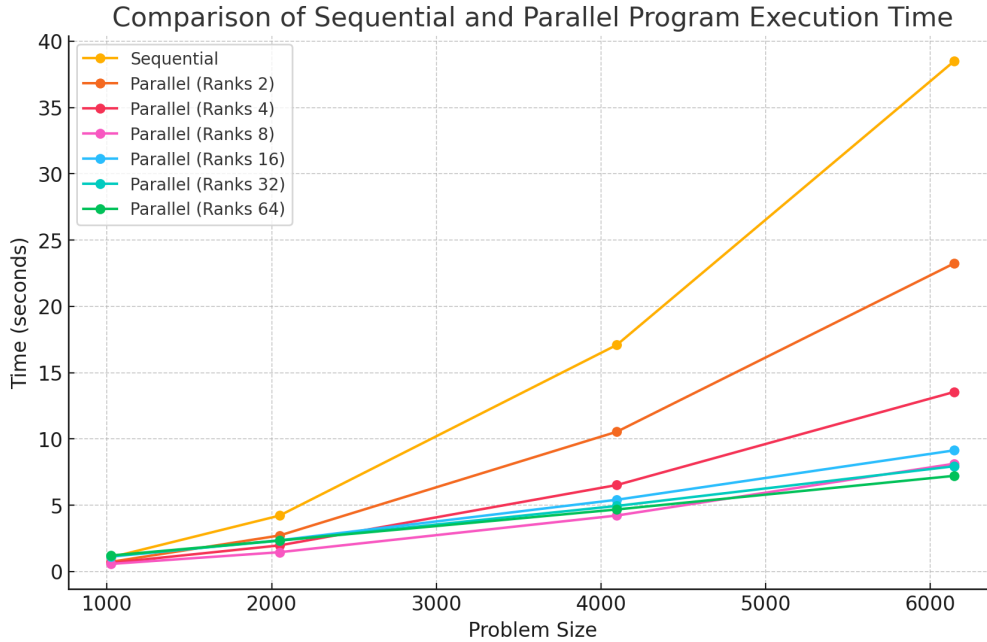
```

```

197
198 Vector createVector(int N) {
199     return malloc(sizeof(value_t) * N);
200 }
201
202 void releaseVector(Vector m) {
203     free(m);
204 }
205
206 void printTemperature(Vector m, int N) {
207     const char *colors = " .-:=+*^X#%@";
208     const int numColors = 12;
209
210     // Boundaries for temperature (for simplicity hard-coded)
211     const value_t max = 273 + 30;
212     const value_t min = 273 + 0;
213
214     // Set the 'render' resolution
215     int W = RESOLUTION;
216
217     // Step size in each dimension
218     int sW = N / W;
219
220     // Room
221     // Left wall
222     printf("X");
223     // Actual room
224     for (int i = 0; i < W; i++) {
225         // Get max temperature in this tile
226         value_t max_t = 0;
227         for (int x = sW * i; x < sW * i + sW; x++) {
228             max_t = (max_t < m[x]) ? m[x] : max_t;
229         }
230         value_t temp = max_t;
231
232         // Pick the 'color'
233         int c = ((temp - min) / (max - min)) * numColors;
234         c = (c >= numColors) ? numColors - 1 : ((c < 0) ? 0 : c);
235
236         // Print the average temperature
237         printf("%c", colors[c]);
238     }
239     // Right wall
240     printf("X");
241 }

```

Listing 3: Parallel propagation



## 2.4 Discuss the effects and implications of your parallelization.

### 1. Performance Improvements

With MPI, the problem is split into parts, and each process (rank) works on one part. This way, the work is shared between cores or nodes, which makes the simulation faster compared to the sequential version.

### 2. Data Communication Overhead

The parallel version needs processes to exchange boundary information using `MPI_Sendrecv`. This adds communication overhead, especially as the number of processes increases. Also, collective functions like `MPI_Bcast` and `MPI_Gather` make all processes wait for each other, which can slow down the program if not managed well.

### 3. Scalability Considerations

Each process handles the same amount of work. But if the problem size does not divide evenly between processes (e.g.,  $N \% \text{numProcs} \neq 0$ ), the program stops. More processes can also increase communication overhead, especially if there is too much communication compared to the actual computation.

### 4. Potential Bottlenecks

Global synchronization using `MPI_Bcast` and `MPI_Gather` can cause delays because all processes must wait for the slowest one to finish its work. This may lead to idle time, reducing efficiency.